

STL Containers

Container basics

An STL container is a collection of objects of the same type (the elements).


- Container *owns* the elements.
 - Creation and destruction is controlled by the container.

Two basic types of containers:

- Sequences
 - User controls the order of elements.
 - vector, list, deque
- Associative containers
 - The container controls the position of elements within it.
 - Elements can be accessed using a *key*.
 - set, multiset, map, multimap

Container concepts

There are three main container concepts:



Container concepts are not as important for generic programming as iterator concepts.

- There are fewer models.
- Containers have important properties that are not described by the basic container concepts.
 - Properties that differentiate one container from another.
 - In contrast, an iterator is almost fully described by the most refined concept it models.
 - More refined container concepts would have just one model.
- Even the basic concepts can be too refined:
 - `boost::array`
- There are almost no generic algorithms taking a container as an argument.
 - insert iterators

However, container concepts standardize basic features.

- Consistent interface makes using them easier.
- Fairly easy to replace one container with another.

The Container concept

Properties shared by all STL containers.

- default constructor
- copy constructor and assignment
 - deep copy
- swap
 - `a.swap(b)` and `swap(a, b)`
 - constant time
- `==`, `!=`
 - content-based equality: equal elements in same order
- order comparisons
 - lexicographic order: first unequal elements determine the order

```
vector<int> a, b;
// a = [1, 2, 3]
// b = [1, 3, 2]
assert(a < b);
```

- `begin()`, `end()`
- `size()`, `empty()`, `max_size()`
- member types
 - `value_type`
 - `reference` (to the value type)
 - `const_reference`
 - `iterator`
 - `const_iterator`
 - `difference_type` (as with iterators)
 - `size_type` (often unsigned type, usually `size_t`)

In addition, a *reversible* container has the properties:

- `rbegin()`, `rend()`
- member types
 - `reverse_iterator`
 - `const_reverse_iterator`

Sequences

Common properties of all sequence containers:

- constructors
 - Fill constructor `Container(n, val)` fills container with `n` copies of `val`.
 - Default fill constructor `Container(n)` fills container with `n` default constructed values.

- Range constructor `Container(i, j)` fills container with the contents of the iterator range `[i,j)`.
- `assign`
 - fill assignment `assign(n, val)`
 - range assignment `assign(i, j)`
 - old elements are assigned to or destroyed
- `insert`
 - `insert(p, val)` inserts `val` just before the position pointed by iterator `p`.
 - `insert(p, n, val)` inserts `n` copies.
 - `insert(p, i, j)` inserts the contents of range `[i,j)`.
- `erase`
 - `erase(p)` erases the element pointed by iterator `p`.
 - `erase(p,q)` erases the range `[p,q)`
 - returns iterator to the position immediately following the erased element(s)
- `clear()` erases all

vector

vector should be used by default as the (sequence) container:

- It is more (space and time) efficient than other STL containers.
- It is more convenient and safer than primitive array.
 - automatic memory management
 - rich interface

Properties of `vector` in addition to sequences:

- `v[i]`, `at(i)`
 - `v.at(i)` checks that $0 \leq i < v.size()$
- `front()`, `back()`
 - return reference to the first and last element (not beyond last)
- `push_back(val)`
 - inserts `val` to the end
- `pop_back()` removes
 - removes the last element and returns it
- `resize`
 - change the number of elements in the vector
 - `resize(n)` makes `n` the size; fills with default values if necessary
 - `resize(n, val)` fills with `val` if necessary
- `capacity()`, `reserve(n)` (see below)

Memory management

- The elements are stored into a contiguous memory area on the heap.
 - `capacity()` is the number of elements that fit into the area.
 - `size()` is the actual number of elements. The remainder of the area is unused (raw memory).
 - `reserve(n)` increases the capacity to `n` without changing the size.
 - The capacity is increased automatically if needed due to insertions.
- Capacity increase may cause copying of all elements.

- A larger memory area is obtained and elements are copied there.
- Capacity increase by an insertion doubles the capacity to achieve *amortized constant time*.
- Capacity never decreases.
 - Memory is not released.
 - But the following gets rid of all extra capacity/memory:

```
vector<int> v;
...
vector<int>(v).swap(v); // copy and swap
```

- Use `&v[0]` to obtain a pointer to the memory area.
 - May be needed as an argument to non-STL functions.

```
vector<char> v(12);
strcpy(&v[0], "hello world");
```

Limitations of vector

- Insertions and deletions in the beginning or in the middle are slow.
 - Requires moving other elements.
 - Prefer `push_back()` and `pop_back()`.
 - Insert or erase many elements at a time by using the range forms of `insert` and `erase`.
- Insertions and deletions *invalidate* all iterators, and pointers and references to the elements.

```
vector<int> v;
...
vector<int> b = v.begin();
v.push_back(x);
find(b, v.end()); // error: b is invalid
```

deque

deque stands for double-ended queue. It is much like vector.


- Differences to vector
 - Insertion and deletion in the beginning in (amortized) constant time.
 - `push_front`, `pop_front`
 - Slower element access and iterators.
 - No `capacity()` or `reserve(n)` but also less need for them.
 - Insertions and deletions to the beginning and end do not invalidate pointers and references to other elements.
 - But iterators may be invalidated.

```
deque<int> d(5,1);
deque<int>::iterator i = d.begin() + 2;
```

```
int* p = &i;
d.push_front(2);
int x = *p; // OK
int y = *i; // error: i may be invalid
```

Memory management

deque stores the elements something like this:



- Element access and iterators are more complicated.
- Fast insertions and deletions to the beginning.
- Handles size changes gracefully: - Capacity increases or decreases one block at a time.
- Memory area is not contiguous.

list

The third standard sequence container is `list`. The underlying data structure is a doubly-linked list:

- No random access.
- Fast insertion and deletion anywhere.
- Insertions and deletions do not invalidate iterators, pointers or references to other elements.

Member functions in addition to sequence:

- `push_front`, `pop_front`
- `push_back`, `pop_back`
- `splice`
 - `c1.splice(i1, c2)` removes all elements from list `c2` and inserts them at position `i1` in list `c1`.
 - `c1.splice(i1, c2, i2)` removes the element pointed by `i2` from `c2` and inserts it at position `i1` in list `c1`.
 - `c1.splice(i1, c2, i2, j2)` removes the range `[i2, j2)` from `c2` and inserts it at position `i1` in list `c1`.
 - In the last two cases, `c1` and `c2` can be the same list.
 - constant time
 - Iterators, pointers and references keep pointing to the same element even if it is moved to a different list.

```
template <class T, class A> // A is allocator
void catenate (list<T,A>& c1, list<T,A>& c2) {
    c1.splice(c1.end(), c2);
}
```

- member versions of STL algorithms
 - `reverse`, `sort`, `merge`, `remove`, `remove_if`, `unique`

- In some cases, like `sort`, the algorithm would not work for lists because it requires random access iterators.
- With the member versions, iterators, pointers and references keep pointing to the same element, unless that element is deleted.

```
list<int> c;
c.push_back(10);
c.push_back(20);
c.push_back(30);
list<int>::iterator i = c.begin();
assert(*i==10); // i points to the first element: 10
c.reverse();
assert(*i==10); // i continues to point to 10
                // which is now to the last element
reverse(c.begin(), c.end());
assert(*i==30); // i still points to the last element
                // which now contains 30
```

string

`string` class was designed before STL, but STL container properties were added to it later. It is similar to vectors; the differences include:

- only `char` as element type
- many additional operations
 - concatenation (`operator+`, `append`)
 - I/O (`<<`, `>>`, `getline`)
 - C-string conversions
 - `substr`, `compare`, `find`, `replace`
 - Many operations can take C-string or substring as an argument.
 - Many of the operations could be replaced with STL algorithms.
- Many implementations have optimizations:
 - reference counting with COW (copy on write)
 - short string optimization

vector<bool>

`vector<bool>` has some special properties.

- Elements are stored as bits in a bit vector.
 - very space-efficient
 - some operations may be slow
- It is not possible to have a pointer or a reference to a bit.
 - `operator[]`, `front()`, `back()`, and iterator's `operator*` do not return a reference but a *proxy* object that behaves almost like a reference but not quite.
 - Taking address of a proxy and assigning it to a reference is not possible.

```
vector<bool> v;
// ...
bool* p = &v[0];    // illegal
bool& r = v.back(); // illegal
```

- Otherwise a proxy can be used on either side of an assignment.

```
bool tmp = v[0];
v[0] = v[1];
v[1] = tmp;
```

- Does not satisfy all container requirements and iterators do not satisfy all requirements of random access iterators.

- flip

- v.flip() flips all bits
- flip one bit: v[1].flip(), v.back().flip(), v.begin()->flip()

Associative containers

The STL standard associative containers (set, multiset, map, multimap) allow access to elements using a **key**:

- For set and multiset element is its own key.
- For map and multimap elements are of type pair<const Key, T>.
 - pair is a standard template class defined as:

```
template <class T, class U>
struct pair {
    T first;
    U second;

    // some constructors
};
```

- set and map contain at most one element for each key.
- multiset and multimap can contain many elements with the same key.

The underlying data structure is a balanced search tree:

- logarithmic access time
- requires order comparisons of keys
- iteration in key order
- Iterators, pointers and references stay valid until the pointed to element is removed.

The order comparison

- operator< by default but can be changed

```
struct my_less {
    bool operator() (int a, int b) { return a < b; }
};
```

```
// all three sets have the same ordering:
set<int> s1; // default: operator< as key order
set<int, std::less<int>> s2;
set<int, my_less> s3;
```

- Two keys are *equivalent* if neither is smaller than the other.
 - `operator==` is not used for comparing keys.
 - Ensures consistency of order and equivalence.
- must be *strict weak ordering*:
 1. *irreflexivity*: $x < x$ is always false.
 2. *transitivity*: $(x < y) \ \&\& \ (y < z)$ implies $x < z$.
 3. *transitivity of equivalence*: if x equals y and y equals z , then x equals z .

```
// NOT strict weak ordering:
// 1 equals 2 and 2 equals 3 but 1 does not equal 3
struct clearly_less {
    bool operator() (int a, int b) { return a < b-1; }
}
```

 - asymmetry: $x < y$ implies $!(y < x)$.
 - Often mentioned as a requirement, but it is implied by 1. and 2.
 - Often called antisymmetry in STL literature, but asymmetry is the correct mathematical term.

Common properties

In addition to properties of the `Container` concept, all associative containers have:

- member types
 - `key_type`
 - `key_compare`
- comparison operators
 - `key_comp()` returns the key comparison operator
 - `value_comp()` returns a comparison operator comparing elements not keys.
- constructors
 - Range constructor `Container(i,j)` fills container with the contents of the range $[i,j)$.
 - (All constructors accept a comparison object as an extra optional argument.)
- `insert`
 - `insert(x)`
 - `insert(i, x)`. Iterator i is a hint pointing to where the search for insertion position should start.
 - Allows `insert_iterator` to operate on associative containers.
 - range insert `insert(i, j)`
 - For `set` and `map` insertions are not done if the key is already there.
- `erase`
 - `erase(k)` erases all elements with key k
 - `erase(i)` erase element pointed to by i
 - range erase `erase(i,j)`
- searching
 - `find(k)` returns iterator to the element with key k or `end()` if no such element
 - `count(k)`
 - `lower_bound(k)` find first element with key not less than k
 - `upper_bound(k)` find first element with key greater than k
 - `equal_range(k)` returns `pair<iterator,iterator>` representing the range of element with key k

set and multiset

- Defined in header file `set`.
- Implement the abstract data structures of set and multiset.
- There are no additional member operations.

```
// count distinct words
set<string> words;
string s;
while (cin >> s) words.insert(s);
cout << words.size() << " distinct words\n";
```

There are no member operations for set intersection, union, etc. However, the following generic algorithms work on any sorted range, including `[s.begin(), s.end())` for a set or multiset `s`:

- `includes`
- `set_intersection`
- `set_union`
- `set_difference`
- `set_symmetric_difference`

```
string str1("abcabcbac");
string str2("abcdabcd");

multiset<char> mset1(str1.begin(), str1.end());
multiset<char> mset2(str2.begin(), str2.end());

multiset<char> result;
set_intersection (mset1.begin(), mset1.end(),
                  mset2.begin(), mset2.end(),
                  inserter(result, result.begin()) );

copy(result.begin(), result.end(), ostream_iterator<char>(cout));
// outputs: "aabbcc"
```

map and multimap

- Defined in header file `map`.
- `multimap` has no additional operations, `map` has one, `operator[]`.
- The elements are pairs, which can make insertion and access slightly awkward.
 - `operator[]` is the most convenient way:

```
map<string, int> days;
days["january"] = 31;
days["february"] = 28;
// ...
days["december"] = 31;
if (leap_year) ++days["february"];

cout << "February has " << days["february"] << " days.\n";
```

- `multimap` does not have `operator[]`. The helper function `make_pair` is useful:

```

multimap<string, string> children;
children.insert(make_pair("Jane","Ann"));
children.insert(make_pair("Jane","Bob"));
children.insert(make_pair("Bob","Xavier"));
// ...

typedef multimap<string, string>::iterator iterator;
pair<iterator,iterator> answer;
answer = children.equal_range("Jane");
cout << "Jane's children:";
for (iterator i = answer.first; i != answer.second; ++i)
    cout << " " << i->second;

```

Container adaptors

The container adaptors `stack`, `queue`, `priority_queue` are containers implemented on top of another container.

They provide a limited set of container operations:

- member types `value_type` and `size_type`, `container_type`
- basic constructors, destructors and assignment
- construction from the underlying container `adaptor(const container&)`
- comparison operators
- `size()`, `empty()`

Stack

`stack` can be implemented on top of `vector`, `deque` or `list`.

- The default is `deque`.

```

// these are equivalent
stack<int> st1;
stack<int, deque<int> > st2;

```

Additional operations:

- constructor `stack(const container&)`
- `push(val)`
- `top()`
- `pop()`

Queue

`queue` can be implemented on top of `deque` or `list`.

- The default is `deque`.

Additional operations:

- `front()`
- `back()`
- `push(val)`

- `pop()`

Priority queue

`priority_queue` can be implemented on top of `deque` or `vector`.

- The default is `vector`.

Uses order comparison operators of elements similar to the associative containers.

```
// these are equivalent
priority_queue<int> pq1;
priority_queue<int, vector<int> > pq2;
priority_queue<int, vector<int>, less<int> > pq3;
```

Additional operations:

- range constructor
- comparison object as an extra optional argument of constructors
- `push(val)`
- `top()` returns the largest element
- `pop()` removes the largest element

There is no method for changing the priority of an element or removing an element that is not the largest.

- Sufficient for some applications like event simulation.
- Not sufficient for others like Dijkstra's algorithm.

Hash tables

The standard has no containers using hashing, but they are a common extension. They are also included in the [Technical Report on C++ Standard Library Extensions](#), commonly known as [TR1](#), an extension of the standard library likely to be included in the next C++ standard:

- `unordered_set`
- `unordered_multiset`
- `unordered_map`
- `unordered_multimap`

More details can be found in the [proposal](#).

Container elements

The type of container elements should always be a model of the `Assignable` concept:

- Normally behaving copy constructor and copy assignment.
- Never store into a container types that are not assignable:
 - references (no assignment)
 - `std::auto_ptr` (abnormal copy behavior)

Some member functions have additional requirements:

- default constructor
 - default fill constructor `Container(n)`
- equality comparisons
 - containers's `operator==`
- order comparisons
 - container's `operator<`
 - associative container with default order comparison

Pointers in containers

Pointers as container elements require special care. There two kinds of pointers:

- Pointers that do not own the object they point to.
 - Example: the same element in multiple containers.
 - for example, different iteration orders
 - One container stores the elements, others store pointers to the elements.
 - Prefer iterators to pointers.
 - They enable container manipulation.
 - Be careful about validity
 - With `deque` use pointers instead of iterators if there are insertions or deletions at the beginning or the end.
 - With `vector` use index if there are insertions or deletions at the end.

- Pointers that own the element they point to.

- Example: polymorphic container:

```
struct animal {
    virtual ~animal() {};
    virtual void eat() =0;
    // ...
}

struct baboon : animal {
    // ...
}

struct lion : animal {
    // ...
}

// ...

vector<animal*> zoo;
zoo.push_back(new baboon);
zoo.push_back(new lion);
zoo[1]->eat();
```

- Such pointers are problematic elements.
 - Containers take care that the destructor of an element is called, when the element is erased (or the container is destroyed).
 - But an owning pointer's destructor does not do what it should: destroy the pointed to object and release the memory.
 - The user of the container must ensure that the pointed to objects are properly destroyed and freed. This is inconvenient and error-prone.

- Achieving exception safety is difficult.
- Better to use a *smart pointer*, whose destructor does the right thing.
 - `auto_ptr` has the right kind of destructor, but unfortunately the wrong kind of copy constructor and assignment.
 - Use `boost::shared_ptr` if possible.

```
typedef boost::shared_ptr<animal> animal_ptr;  
vector<animal_ptr> zoo;  
  
animal_ptr p(new baboon);  
zoo.push_back(p);  
p.reset(new lion);  
zoo.push_back(p);  
zoo[1]->eat();
```

Exception safety

Exceptions are a common error reporting mechanism in C++. The elements of STL containers are allowed to throw exceptions (except in destructors). In particular, if a copy of an element fails and throws an exception during a container operation, one of the following guarantees are provided:

- *Strong guarantee*: The container operation is cancelled, and the container's state is as if the operation was never called.
 - Most list operations.
 - All single element operations on lists and associative containers.
 - push and pop operations on vector and deque.
- *Basic guarantee*: There are no memory leaks, and the container is in a consistent state. In particular, the container can be destroyed safely.

These guarantees make it possible to recover from an exception without memory leaks.